


# Evaluating Apple's MLX Framework for Machine Learning


A Comparative Study Using Unsupervised  
Clustering Algorithms

Sharan Deepak Thakur  
GH1031360



# Agenda

1. Introduction
2. Literature Review
3. The Research Gap
4. Research Questions
5. Methodology & Conceptual Framework
6. Experimental Methodology
7. Challenges
8. Results
9. Conclusion



Generated using AI

# Introduction





Fig: Machine Learning Problems  
Dan Shewan WordStream, 2023



# The Evolution of Machine Learning Hardware

1

## GPU-Centric Era

Traditional deep learning frameworks optimized for NVIDIA CUDA architecture with discrete GPU memory.

2

## Apple Silicon Shift

1. M-series processors  
2. Unified memory architecture integrating CPU, GPU, ANE with shared pools up to 192GB.

3

## MLX Framework

Purpose-built array framework designed specifically to leverage Apple Silicon's unique architecture from the ground up.

# Literature Review

Paper Author	Relevance
<i>Awni Hannun et al., 2025</i>	Defines the MLX framework's architecture, optimizations, and rationale for Apple Silicon benchmarking.
<i>Indranil Bose &amp; Xi Chen, 2009</i>	It formalizes both hybridization strategies (cluster → classifier and cluster label as feature).
<i>Hoang Tran, Ngoc Le, Van-Ho Nguyen, 2023</i>	Uses K-Means for segmentation on financial data (credit card/banking).

Table: Relevant Papers

## KMeans Clustering:

Hartigan and Wong. 1979, Most widely deployed unsupervised algorithm for customer segmentation, financial analytics, and exploratory analysis

## MLX Design Philosophy:

Hannun et al. 2025 designed MLX with lazy evaluation and unified memory utilization

## Apple Silicon Architecture:

Feng et al. 2025 reveals comprehensive profiling unified memory eliminates CPU-GPU data transfer overhead

# The Research Gap

Current literature overwhelmingly focuses on MLX performance for deep learning specifically LLMs and transformers. However, **a critical gap exists for traditional machine learning algorithms.**

Algorithmic Diversity

Production Relevance



## Research Questions

1

### O1. Performance Comparison

RQ1. *How does MLX's computational performance for KMeans clustering compare to scikit-learn, NumPy, and PyTorch on Apple Silicon?*

2

### O2. Clustering Quality

RQ2. *Does MLX produce equivalent results measured by Adjusted Rand Index, Normalized Mutual Information, and Silhouette Score?*

3

### O3. API Usability

RQ3. *What are the primary implementation challenges when developing traditional ML algorithms in MLX given its API constraints?*

4

### O4. Practical Implications

RQ4. *What actionable recommendations should practitioners consider when evaluating MLX for unsupervised learning workflows?*

# Methodology & Conceptual Framework




Fig: Proposed Methodology for experiment.

Made with [mermaidchart.com](https://mermaidchart.com)




Fig: Conceptual Framework

## Hypotheses

- **H1:** Clustering Quality → Label
  - Is the clustering quality aligned with the churn labels of the labelled data?
- **H2:** Memory Efficiency
  - Does MLX use less memory than other frameworks.

# Experimental Methodology

01

## Multi-Framework Implementations

Parallel implementations of KMeans across MLX, scikit-learn, NumPy, and PyTorch using identical algorithmic logic and parameters.

02

## Rigorous Benchmarking Protocol

Experimental rounds following Bahrampour et al. from 2016's framework comparison methodology.

03

## Controlled Environment

All experiments executed on identical Apple Silicon hardware, Python 3.11, with standardized data preprocessing and feature scaling between 0-1.

04

## Comprehensive Metrics

Clustering quality evaluation using


$$ARI = \frac{\sum_{ij} \binom{n_{ij}}{2} - [\sum_i \binom{a_i}{2} \sum_j \binom{b_j}{2}] / \binom{n}{2}}{\frac{1}{2} [\sum_i \binom{a_i}{2} + \sum_j \binom{b_j}{2}] - [\sum_i \binom{a_i}{2} \sum_j \binom{b_j}{2}] / \binom{n}{2}}$$

$$NMI(U, V) = \frac{2 \cdot I(U; V)}{H(U) + H(V)}$$

$$s(i) = \frac{b(i) - a(i)}{\max(a(i), b(i))} \quad a(i) = \text{mean intra-cluster distance}, \quad b(i) = \text{mean nearest-cluster distance}$$

- Computational performance measured via wall-clock time with millisecond precision.

*ARI:* L. Hubert et al, 1985; *NMI:* Scikit-Learn, *Silhouette Score:* Peter. J.R., 1987



# MLX Implementation Challenges

## Challenge:

No Boolean Indexing

**Problem:** MLX lacks direct support for operations like `X[labels==k]`.

**Solution:** Implemented mask-based weighted operations to achieve equivalent functionality.

## Challenge:

Limited Random API

**Problem:** No equivalent to NumPy's `np.random.choice()`.

**Solution:** Used `mx.random.permutation()` followed by slicing for centroid initialization.

## Challenge:

Type System Differences


**Problem:** MLX arrays require careful type checking and explicit conversions.

**Solution:** Implemented comprehensive input validation and type conversion logic throughout.

These API differences reveal MLX's current difference in their array operations due to lazy computation.

# Preliminary Results: MNIST Dataset

**Configuration:** 7,500 samples with 784 features, k=10 clusters, 20 experimental rounds with seeds 44-63



## Clustering Quality

All frameworks achieve equivalent quality metrics within 5%. MLX achieves the **highest ARI (0.3777) and Silhouette (0.0636)**, demonstrating superior clustering results.

**ARI (Adjusted Rand Index)** measures similarity between true and predicted clusters, ranging from -1 to 1, where 1 indicates perfect matching.


**Silhouette Score** measures how similar an object is to its own cluster compared to other clusters, ranging from -1 to 1, with higher values indicating better-defined clusters.

## Performance Ranking

**scikit-learn** (0.16s) > PyTorch (0.61s) > MLX (1.10s) > NumPy (2.78s). MLX achieves **60% speedup over NumPy**.

# Final Results: Credit Card Dataset

**Configuration:** 30,000 samples with 9 features, k=2 clusters, 500 KMeans rounds with seed 44



## Clustering Quality

- H1 Refuted - Clustering Does Not Align with Churn**
- Silhouette Score:** ~0.52 for all implementations
- No correlation** with the hidden churn labels (avg. **ARI**  $\approx$  -0.03).
  
- H2 Supported – MLX is Highly Memory Efficient**
- MLX used less memory than sklearn and numpy.
  
- Performance Ranking is Data-Dependent**

## Performance Ranking

NumPy (0.0245s) > MLX (0.0401s) > sklearn (0.1805s) > PyTorch (0.1955s)

# Implications & Recommendations

## For Research Community

- MLX is a viable, lightweight, and memory-efficient framework for classical ML algorithms, not just LLMs.

## For Practitioners

- K-Means should **not** be used as a proxy for churn prediction on *this* dataset. The discovered segments, while statistically valid, are not relevant to *this* specific business problem.
- For small-to-medium datasets where **memory efficiency is critical** (e.g., on-device or shared-resource environments), **MLX is a superior choice** over sklearn.

## For Future Work

- Investigate why the clusters failed to capture churn (e.g., feature engineering, feature selection)
- Benchmark scaling.

This research fills a critical gap in machine learning systems literature by providing empirical evidence that MLX can effectively support traditional algorithms, not just deep learning, on Apple Silicon hardware.

# Thank You

